CS294-248 Special Topics in Database Theory Unit 6: Constraints, Incomplete and Probabilistic Databases

Dan Suciu

University of Washington

Outline

- Today: Generalized Constraints, Semantics Optimization.

- Thursday: Repairs, Incomplete Databases

Constraints

A constraint is an assertion on the database D that must always hold.

How does this differ from <u>invariants</u> in programs?

Constraints

A constraint is an assertion on the database **D** that must always hold.

How does this differ from <u>invariants</u> in programs?

Constraints: we check them at runtime (this may be costly)

Invariants: we prove them offline, do not check at runtime.

Recap
○●

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○○○

Applications of Constraints

- Enforce database consistency.
    Most common constraint in practice:

    "Please type in your phone number using $XXX - XXX - XXXX$";

- Database normalization.

- Semantic optimization: given query Q find a "better" query Q' s.t. $Q \equiv Q'$ on databases satisfying the constraints.

- Database repair: if $\boldsymbol{D} \not\models \Sigma$, delete/insert tuples s.t. $\boldsymbol{D'} \models \Sigma$.

- Consistent query answering: given query Q return only those answers that are present in $Q(\boldsymbol{D'})$ for all repairs $\boldsymbol{D'}$.

Recap
○○

Classical Database Constraints
●○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○○○

Classical Database Constraints

Classical Database Constraints

- Functional Dependencies (FD).

- Multivalued Dependencies (MVD).

- Join Dependencides (JD).

- Inclusion Dependencies (IND).

Recap
○○

Classical Database Constraints
○○●○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○○○

Functional Dependency

Notation:
$$\boxed{\boldsymbol{U} \rightarrow \boldsymbol{V}}$$

Semantics: $R^D \models \boldsymbol{U} \rightarrow \boldsymbol{V}$ if:

$$\forall \boldsymbol{u}, \boldsymbol{v}_1, \boldsymbol{w}_1, \boldsymbol{v}_2, \boldsymbol{w}_2 (R(\boldsymbol{u}, \boldsymbol{v}_1, \boldsymbol{w}_1) \wedge R(\boldsymbol{u}, \boldsymbol{v}_2, \boldsymbol{w}_2) \Rightarrow \boldsymbol{v}_1 = \boldsymbol{v}_2)$$

Consequence

Lossless decomposition: $R(\boldsymbol{U}, \boldsymbol{V}, \boldsymbol{W}) = R_1(\boldsymbol{U}, \boldsymbol{V}) \bowtie R_2(\boldsymbol{U}, \boldsymbol{W})$.

The implication problem: axiomatizable (Armstrong), decidable in PTIME.

Multivalued Dependency

Notation: given a partition all attribute $\boldsymbol{X} = \boldsymbol{U} \cup \boldsymbol{V} \cup \boldsymbol{W}$:

$$\boxed{\boldsymbol{U} \twoheadrightarrow \boldsymbol{V}; \boldsymbol{W}}$$

Semantics: $R^D \models \boldsymbol{U} \twoheadrightarrow \boldsymbol{V}; \boldsymbol{W}$ if:

$$\forall \boldsymbol{u}, \boldsymbol{v}_1, \boldsymbol{w}_1, \boldsymbol{v}_2, \boldsymbol{w}_2(R(\boldsymbol{u}, \boldsymbol{v}_1, \boldsymbol{w}_1) \wedge R(\boldsymbol{u}, \boldsymbol{v}_2, \boldsymbol{w}_2) \Rightarrow R(\boldsymbol{u}, \boldsymbol{v}_1, \boldsymbol{w}_2))$$

Equivalent Definition

Lossless decomposition: $R(\boldsymbol{U}, \boldsymbol{V}, \boldsymbol{W}) = R_1(\boldsymbol{U}, \boldsymbol{V}) \bowtie R_2(\boldsymbol{U}, \boldsymbol{W})$.

The implication problem for FD+MVD: axiomatizable, decidable.

Recap
○○

Classical Database Constraints
○○○○●○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○○○

Join Dependencies

Notation: given a cover of all attributes $\boldsymbol{X} = \boldsymbol{U}_1 \cup \ldots \cup \boldsymbol{U}_k$:

$$\bowtie (\boldsymbol{U}_1, \boldsymbol{U}_2, \ldots, \boldsymbol{U}_k)$$

Join Dependencies

Notation: given a cover of all attributes $\boldsymbol{X} = \boldsymbol{U}_1 \cup \ldots \cup \boldsymbol{U}_k$:

$$\bowtie (\boldsymbol{U}_1, \boldsymbol{U}_2, \ldots, \boldsymbol{U}_k)$$

Semantics by example. $R^D(X, Y, Z) \models \bowtie (XY, YZ, XZ)$ if R^D satisfies

$$\forall x, x', y, y', z, z'(R(x, y, z') \wedge R(x', y, z) \wedge R(x, y', z)) \Rightarrow R(x, y, z)$$

Join Dependencies

Notation: given a cover of all attributes $\boldsymbol{X} = \boldsymbol{U}_1 \cup \ldots \cup \boldsymbol{U}_k$:

$$\bowtie (\boldsymbol{U}_1, \boldsymbol{U}_2, \ldots, \boldsymbol{U}_k)$$

Semantics by example. $R^D(X, Y, Z) \models \bowtie (XY, YZ, XZ)$ if R^D satisfies

$$\forall x, x', y, y, z, z'(R(x, y, z') \land R(x', y, z) \land R(x, y', z)) \Rightarrow R(x, y, z)$$

Equivalently: $R^D \models \bowtie (\boldsymbol{U}_1, \ldots, \boldsymbol{U}_k)$ if:

Definition of JD

Lossless decomposition: $R(\boldsymbol{X}) = R_1(\boldsymbol{U}_1) \bowtie \cdots \bowtie R_k(\boldsymbol{U}_k)$

Join Dependencies

Notation: given a cover of all attributes $\boldsymbol{X} = \boldsymbol{U}_1 \cup \ldots \cup \boldsymbol{U}_k$:

$$\bowtie (\boldsymbol{U}_1, \boldsymbol{U}_2, \ldots, \boldsymbol{U}_k)$$

Semantics by example. $R^D(X, Y, Z) \models \bowtie (XY, YZ, XZ)$ if R^D satisfies

$$\forall x, x', y, y, z, z'(R(x, y, z') \wedge R(x', y, z) \wedge R(x, y', z)) \Rightarrow R(x, y, z)$$

Equivalently: $R^D \models \bowtie (\boldsymbol{U}_1, \ldots, \boldsymbol{U}_k)$ if:

Definition of JD

Lossless decomposition: $R(\boldsymbol{X}) = R_1(\boldsymbol{U}_1) \bowtie \cdots \bowtie R_k(\boldsymbol{U}_k)$

JD implication problem not axiomatizable [Abiteboul et al., 1995, pp.171].

FD+JD implication problem is decidable (later).

Inclusion Dependencies

Notation: relation schemas $R(\boldsymbol{X}), S(\boldsymbol{Y})$, $\boldsymbol{U} \subseteq \boldsymbol{X}$, $\boldsymbol{V} \subseteq \boldsymbol{Y}$, $|\boldsymbol{U}| = |\boldsymbol{V}|$:

$$\boxed{R[\boldsymbol{U}] \subseteq S[\boldsymbol{V}]}$$

Inclusion Dependencies

Notation: relation schemas $R(\boldsymbol{X}), S(\boldsymbol{Y})$, $\boldsymbol{U} \subseteq \boldsymbol{X}$, $\boldsymbol{V} \subseteq \boldsymbol{Y}$, $|\boldsymbol{U}| = |\boldsymbol{V}|$:

$$\boxed{R[\boldsymbol{U}] \subseteq S[\boldsymbol{V}]}$$

Semantics: (what you expect, but watch the FO sentence):

$$\boxed{\forall \boldsymbol{u} \forall \boldsymbol{r}(R(\boldsymbol{u}, \boldsymbol{r}) \Rightarrow \exists \boldsymbol{s} S(\boldsymbol{u}, \boldsymbol{s}))}$$

Recap
○○

Classical Database Constraints
○○○○○●○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○○○

Inclusion Dependencies

Notation: relation schemas $R(\boldsymbol{X}), S(\boldsymbol{Y})$, $\boldsymbol{U} \subseteq \boldsymbol{X}$, $\boldsymbol{V} \subseteq \boldsymbol{Y}$, $|\boldsymbol{U}| = |\boldsymbol{V}|$:

$$\boxed{R[\boldsymbol{U}] \subseteq S[\boldsymbol{V}]}$$

Semantics: (what you expect, but watch the FO sentence):

$$\boxed{\forall \boldsymbol{u} \forall \boldsymbol{r}(R(\boldsymbol{u}, \boldsymbol{r}) \Rightarrow \exists \boldsymbol{s} S(\boldsymbol{u}, \boldsymbol{s}))}$$

[Abiteboul et al., 1995, pp.171-202]:

- IND is axiomatizable (3 simple axioms).
- The implication problem for IND is PSPACE complete.
- The implication problem for FD+IND is undecidable.

Discussion

FDs, MVDs, JDs, INDs, ..., why so many kinds?

It turns out that all can be captured by a single formalism:

Generalized Dependencies

Generalized Dependencies

Generalized Dependency

Relational schema: $R_1, R_2, \ldots$

A Generalized Dependency is a statement of one of these two forms:

Tuple-Generating Dependency (TGD):

$$\forall \boldsymbol{x} \, (A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$$

Generalized Dependency

Relational schema: $R_1, R_2, \ldots$

A Generalized Dependency is a statement of one of these two forms:

Tuple-Generating Dependency (TGD):

$$\forall \boldsymbol{x}\,(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$$

The TGD is full if there is no $\exists \boldsymbol{y}$

Generalized Dependency

Relational schema: $R_1, R_2, \ldots$

A Generalized Dependency is a statement of one of these two forms:

Tuple-Generating Dependency (TGD):

$$\forall \mathbf{x}\,(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \mathbf{y}(B_1 \wedge \cdots \wedge B_k))$$

The TGD is full if there is no $\exists \mathbf{y}$

Equality-Generating Dependency (EGD):

$$\forall \mathbf{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow x_i = x_j)$$

Generalized Dependency

Relational schema: $R_1, R_2, \ldots$

A Generalized Dependency is a statement of one of these two forms:

Tuple-Generating Dependency (TGD):

$$\boxed{\forall \boldsymbol{x}\,(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))}$$

The TGD is full if there is no $\exists \boldsymbol{y}$

Equality-Generating Dependency (EGD):

$$\boxed{\forall \boldsymbol{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow x_i = x_j)}$$

Examples

FD: $\forall u \forall x_1 \forall x_2 (R(u, x_1) \wedge R(u, x_2) \Rightarrow x_1 = x_2)$.

Generalized Dependency

Relational schema: $R_1, R_2, \ldots$

A Generalized Dependency is a statement of one of these two forms:

Tuple-Generating Dependency (TGD):

$$\forall \boldsymbol{x} \, (A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$$

The TGD is full if there is no $\exists \boldsymbol{y}$

Equality-Generating Dependency (EGD):

$$\forall \boldsymbol{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow x_i = x_j)$$

Examples

FD: $\forall u \forall x_1 \forall x_2 (R(u, x_1) \wedge R(u, x_2) \Rightarrow x_1 = x_2)$.

MVD: $\forall u, v_1, w_1, v_2, v_2 (R(u, v_1, w_1) \wedge R(u, v_2, w_2) \Rightarrow R(u, v_1, w_2))$.

Generalized Dependency

Relational schema: $R_1, R_2, \ldots$

A Generalized Dependency is a statement of one of these two forms:

Tuple-Generating Dependency (TGD):

$$\forall \boldsymbol{x} \, (A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$$

The TGD is full if there is no $\exists \boldsymbol{y}$

Equality-Generating Dependency (EGD):

$$\forall \boldsymbol{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow x_i = x_j)$$

Examples

FD: $\forall u \forall x_1 \forall x_2(R(u, x_1) \wedge R(u, x_2) \Rightarrow x_1 = x_2)$.

MVD: $\forall u, v_1, w_1, v_2, v_2(R(u, v_1, w_1) \wedge R(u, v_2, w_2) \Rightarrow R(u, v_1, w_2))$.

IND: $\forall x \forall x'(R(x, x') \Rightarrow \exists y S(x, y'))$.

Dissecting Generalized Dependencies

$\forall \boldsymbol{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$

- Need $\exists$ on the right, but not on the left:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$

Dissecting Generalized Dependencies

$\forall \boldsymbol{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$

- Need $\exists$ on the right, but not on the left:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ equivalent to $\forall x \forall y(R(x, y) \Rightarrow \exists z S(x, z))$

Dissecting Generalized Dependencies

$\forall \mathbf{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \mathbf{y}(B_1 \wedge \cdots \wedge B_k))$

- Need $\exists$ on the right, but not on the left:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ equivalent to $\forall x \forall y(R(x, y) \Rightarrow \exists z S(x, z))$

- When $\exists$ is missing, then we can split the RHS:
  $\forall x \forall y(R(x, y) \Rightarrow S(x) \wedge (x = y))$

Dissecting Generalized Dependencies

$\forall \boldsymbol{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \boldsymbol{y}(B_1 \wedge \cdots \wedge B_k))$

- Need $\exists$ on the right, but not on the left:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ equivalent to $\forall x \forall y(R(x, y) \Rightarrow \exists z S(x, z))$

- When $\exists$ is missing, then we can split the RHS:
  $\forall x \forall y(R(x, y) \Rightarrow S(x) \wedge (x = y))$ equivalent to two GDs:
  $\forall x \forall y(R(x, y) \Rightarrow S(x))$ $\forall x \forall y(R(x, y) \Rightarrow (x = y))$

Dissecting Generalized Dependencies

$\forall \mathbf{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \mathbf{y}(B_1 \wedge \cdots \wedge B_k))$

- Need $\exists$ on the right, but not on the left:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ equivalent to $\forall x \forall y(R(x, y) \Rightarrow \exists z S(x, z))$

- When $\exists$ is missing, then we can split the RHS:
  $\forall x \forall y(R(x, y) \Rightarrow S(x) \wedge (x = y))$ equivalent to two GDs:
  $\forall x \forall y(R(x, y) \Rightarrow S(x))$ $\forall x \forall y(R(x, y) \Rightarrow (x = y))$

- A GD is equivalent to a query containment assertion:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ is equivalent to:
  $Q_1 \subseteq Q_2$ where $Q_1(x) \stackrel{\text{def}}{=} \exists y R(x, y),\ Q_2(x) \stackrel{\text{def}}{=} \exists z S(x, z)$.

Recap
00

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
00000000000

Dissecting Generalized Dependencies

$\forall \mathbf{x}(A_1 \wedge \ldots \wedge A_m \Rightarrow \exists \mathbf{y}(B_1 \wedge \cdots \wedge B_k))$

- Need $\exists$ on the right, but not on the left:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ equivalent to $\forall x \forall y(R(x, y) \Rightarrow \exists z S(x, z))$

- When $\exists$ is missing, then we can split the RHS:
  $\forall x \forall y(R(x, y) \Rightarrow S(x) \wedge (x = y))$ equivalent to two GDs:
  $\forall x \forall y(R(x, y) \Rightarrow S(x))$ $\forall x \forall y(R(x, y) \Rightarrow (x = y))$

- A GD is equivalent to a query containment assertion:
  $\forall x(\exists y R(x, y) \Rightarrow \exists z S(x, z))$ is equivalent to:
  $Q_1 \subseteq Q_2$ where $Q_1(x) \stackrel{\text{def}}{=} \exists y R(x, y)$, $Q_2(x) \stackrel{\text{def}}{=} \exists z S(x, z)$.

- To check $\mathbf{D} \models \sigma$, compute $Q_1(\mathbf{D}), Q_2(\mathbf{D})$; in PTIME.

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○●

Semantics Optimization
○○○○○○○○○○○

Discussion

- GDs are a fragment of FO, powerful enough to capture classical constraints, yet weak enough to be useful.

- Next: we show their utility in semantics optimization.

Recap
Classical Database Constraints
Generalized Dependencies
Semantics Optimization
oo
ooooooo
oooo
●ooooooooooo

Semantic Query Optimization

Overview

Semantics Query Optimization means query optimization that uses the database constraints Σ

Replace a query Q by Q' such that $Q(\boldsymbol{D}) = Q'(\boldsymbol{D})$ for every database instance $\boldsymbol{D}$ that satisfies the constraints.

We write $\Sigma \models Q \equiv Q'$

Note that, in general, $Q \not\equiv Q'$.

Semantic optimization is an old idea
[King, 1981, Chakravarthy et al., 1990].

Example

$$Q_1(z) = R(x, 55) \land R(x, y) \land S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? $Q_2 \subseteq Q_1$?

Example

$$Q_1(z) = R(x, 55) \wedge R(x, y) \wedge S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$?

Recap
oo

Classical Database Constraints
ooooooo

Generalized Dependencies
oooo

Semantics Optimization
oooooooooooo

Example

$$Q_1(z) = R(x, 55) \wedge R(x, y) \wedge S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO

(In class: show canonical database refuting these containments)

Recap
00

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
000000000000

Example

$$Q_1(z) = R(x, 55) \land R(x, y) \land S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○●○○○○○○○○○○

Example

$$Q_1(z) = R(x, 55) \land R(x, y) \land S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?
$R.x$ is a key: $$\sigma_1 : \forall x, y, w(R(x, w) \land R(x, y) \Rightarrow (w = y))$$

Example

$$Q_1(z) = R(x, 55) \land R(x, y) \land S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?
$R.x$ is a key: $\sigma_1 : \forall x, y, w(R(x, w) \land R(x, y) \Rightarrow (w = y))$
Then $Q_1(z) \equiv R(x, 55) \land R(x, 55) \land S(55, z) \equiv R(x, 55) \land S(55, z)$

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○●○○○○○○○○○

Example

$$Q_1(z) = R(x, 55) \wedge R(x, y) \wedge S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?
$R.x$ is a key: $\sigma_1 : \forall x, y, w(R(x, w) \wedge R(x, y) \Rightarrow (w = y))$
Then $Q_1(z) \equiv R(x, 55) \wedge R(x, 55) \wedge S(55, z) \equiv R(x, 55) \wedge S(55, z)$

What constraint implies $Q_2 \subseteq Q_1$?

Example

$$Q_1(z) = R(x, 55) \land R(x, y) \land S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?
$R.x$ is a key: $\sigma_1 : \forall x, y, w(R(x, w) \land R(x, y) \Rightarrow (w = y))$
Then $Q_1(z) \equiv R(x, 55) \land R(x, 55) \land S(55, z) \equiv R(x, 55) \land S(55, z)$

What constraint implies $Q_2 \subseteq Q_1$?
IND: $\sigma_2 : \forall y, z(S(y, z) \Rightarrow \exists x R(x, y))$.

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○●○○○○○○○○○

Example

$$Q_1(z) = R(x, 55) \wedge R(x, y) \wedge S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $\quad Q_1 \subseteq Q_2$? NO $\quad Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?
$R.x$ is a key: $\quad \sigma_1 : \forall x, y, w(R(x, w) \wedge R(x, y) \Rightarrow (w = y))$
Then $Q_1(z) \equiv R(x, 55) \wedge R(x, 55) \wedge S(55, z) \equiv R(x, 55) \wedge S(55, z)$

What constraint implies $Q_2 \subseteq Q_1$?
IND: $\quad \sigma_2 : \forall y, z(S(y, z) \Rightarrow \exists x R(x, y))$. $\quad$ Then:
$Q_2(z) \equiv S(y, z) \wedge y = 55 \equiv R(x, y) \wedge S(y, z) \wedge y = 55 \equiv R(x, 55) \wedge S(55, z)$

Recap
oo

Classical Database Constraints
ooooooo

Generalized Dependencies
oooo

Semantics Optimization
oooooooooooo

Example

$$Q_1(z) = R(x, 55) \land R(x, y) \land S(y, z)$$

$$Q_2(z) = S(55, z)$$

Which of the following hold? $Q_1 \subseteq Q_2$? NO $Q_2 \subseteq Q_1$? NO
(In class: show canonical database refuting these containments)

What constraint implies $Q_1 \subseteq Q_2$?
$R.x$ is a key: $\sigma_1 : \forall x, y, w(R(x, w) \land R(x, y) \Rightarrow (w = y))$
Then $Q_1(z) \equiv R(x, 55) \land R(x, 55) \land S(55, z) \equiv R(x, 55) \land S(55, z)$

What constraint implies $Q_2 \subseteq Q_1$?
IND: $\sigma_2 : \forall y, z(S(y, z) \Rightarrow \exists x R(x, y))$.   Then:
$Q_2(z) \equiv S(y, z) \land y = 55 \equiv R(x, y) \land S(y, z) \land y = 55 \equiv R(x, 55) \land S(55, z)$

Assume the database satisfies σ_1, σ_2. Then we can optimize Q_1 to Q_2

Recap
00

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
000●00000000

The Chase: Overview

- The Chase takes a query Q and a GD σ and creates a new query Q_1 by "applying" σ to Q.

- The important semantics property of the chase is: $\sigma \models Q \equiv Q_1$.

- By repeatedly applying the chase we obtain a sequence $Q, Q_1, Q_2, \ldots$

- To check $\Sigma \models Q \equiv Q'$ it suffices to find a chase sequence from Q to Q_m, and one from Q' to Q'_n, then prove $Q_m \equiv Q'_n$ (unconditioned).

Recap
oo

Classical Database Constraints
ooooooo

Generalized Dependencies
oooo

Semantics Optimization
ooooo●oooooo

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \rightarrow Q$, we write $Q \overset{\sigma,\theta}{\rightarrow} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j/x_i]$.

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \to Q$, we write $Q \overset{\sigma, \theta}{\to} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j/x_i]$.

Example $\boxed{Q(x) = R(x, y) \wedge A(y) \wedge R(x, z) \wedge B(z)}$

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○●○○○○○○

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \to Q$, we write $Q \overset{\sigma, \theta}{\to} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j/x_i]$.

Example $\boxed{Q(x) = R(x, y) \wedge A(y) \wedge R(x, z) \wedge B(z)}$

$$\sigma_1 = \forall u \forall v \forall w(R(u, v) \wedge R(u, w) \Rightarrow (v = w))$$
$$\sigma_2 = \forall u \forall v(R(u, v) \Rightarrow \exists w S(u, w))$$

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○●○○○○○○○

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \to Q$, we write $Q \overset{\sigma, \theta}{\to} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j / x_i]$.

Example $\boxed{Q(x) = R(x, y) \wedge A(y) \wedge R(x, z) \wedge B(z)}$

$\qquad \sigma_1 = \forall u \forall v \forall w(R(u, v) \wedge R(u, w) \Rightarrow (v = w))$

$\qquad \sigma_2 = \forall u \forall v(R(u, v) \Rightarrow \exists w S(u, w))$

Chase Q with σ_1, $\theta_1 : (u, v, w) \mapsto (x, y, z)$.

$$Q \overset{\sigma_1, \theta_1}{\to} ?$$

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \rightarrow Q$, we write $Q \overset{\sigma, \theta}{\rightarrow} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j/x_i]$.

Example $\boxed{Q(x) = R(x,y) \wedge A(y) \wedge R(x,z) \wedge B(z)}$

$\sigma_1 = \forall u \forall v \forall w (R(u,v) \wedge R(u,w) \Rightarrow (v = w))$

$\sigma_2 = \forall u \forall v (R(u,v) \Rightarrow \exists w S(u,w))$

Chase Q with σ_1, $\theta_1 : (u,v,w) \mapsto (x,y,z)$.

$$Q \overset{\sigma_1, \theta_1}{\rightarrow} \boxed{R(x,y) \wedge A(y) \wedge B(y)}$$

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○●○○○○○○

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \to Q$, we write $Q \overset{\sigma, \theta}{\to} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j/x_i]$.

Example $\boxed{Q(x) = R(x, y) \wedge A(y) \wedge R(x, z) \wedge B(z)}$

$$\sigma_1 = \forall u \forall v \forall w (R(u, v) \wedge R(u, w) \Rightarrow (v = w))$$
$$\sigma_2 = \forall u \forall v (R(u, v) \Rightarrow \exists w S(u, w))$$

Chase Q with σ_1, $\theta_1 : (u, v, w) \mapsto (x, y, z)$.

$$Q \overset{\sigma_1, \theta_1}{\to} \boxed{R(x, y) \wedge A(y) \wedge B(y)}$$

Chase the result with σ_2, $\theta_2 : (u, v) \mapsto (x, y)$.

$$Q' \overset{\sigma_2, \theta_2}{\to} ?$$

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○●○○○○○○

Definition of the Chase

Let σ be $\forall \boldsymbol{x}(A \Rightarrow C)$ where A is a conjunction of atoms, Q be a CQ.

Definition (The Chase)

For a homomorphism $\theta : A \to Q$, we write $Q \overset{\sigma, \theta}{\to} Q'$ where Q' is:

- If σ is a TGD with $C \equiv \exists \boldsymbol{y} B$, then $Q' \overset{\text{def}}{=} Q \wedge \theta(B)$.

- If σ is an EGD with $C \equiv (x_i = x_j)$, then $Q' \overset{\text{def}}{=} Q[x_j/x_i]$.

Example $\boxed{Q(x) = R(x, y) \wedge A(y) \wedge R(x, z) \wedge B(z)}$

$$\sigma_1 = \forall u \forall v \forall w (R(u, v) \wedge R(u, w) \Rightarrow (v = w))$$
$$\sigma_2 = \forall u \forall v (R(u, v) \Rightarrow \exists w S(u, w))$$

Chase Q with σ_1, $\theta_1 : (u, v, w) \mapsto (x, y, z)$.

$$Q \overset{\sigma_1, \theta_1}{\to} \boxed{R(x, y) \wedge A(y) \wedge B(y)}$$

Chase the result with σ_2, $\theta_2 : (u, v) \mapsto (x, y)$.

$$Q' \overset{\sigma_2, \theta_2}{\to} \boxed{R(x, y) \wedge A(y) \wedge B(y) \wedge S(x, w)}$$

When the Chase Goes Wrong

- Given a set Σ of GD, we can repeatedly apply the chase:
  $Q \overset{\sigma_1,\theta_1}{\to} Q_1 \overset{\sigma_2,\theta_2}{\to} Q_2 \cdots$

---

[1]The book doesn't consider constants; need to add this to allow constants.

When the Chase Goes Wrong

- Given a set Σ of GD, we can repeatedly apply the chase:
  $Q \overset{\sigma_1, \theta_1}{\to} Q_1 \overset{\sigma_2, \theta_2}{\to} Q_2 \cdots$

- In general, this may not terminate:
  $\sigma = \forall x \forall y (R(x, y) \to \exists z R(y, z))$ $Q() = R(u_0, u_1)$
  $R(u_0, u_1) \to (R(u_0, u_1) \wedge R(u_1, u_2)) \to (R(u_0, u_1) \wedge R(u_1, u_2) \wedge R(u_2, u_3)) \to \cdots$

---

[1]The book doesn't consider constants; need to add this to allow constants.

When the Chase Goes Wrong

- Given a set Σ of GD, we can repeatedly apply the chase:
  $Q \overset{\sigma_1, \theta_1}{\to} Q_1 \overset{\sigma_2, \theta_2}{\to} Q_2 \cdots$

- In general, this may not terminate:
  $\sigma = \forall x \forall y (R(x, y) \to \exists z R(y, z))$ $\qquad\qquad$ $Q() = R(u_0, u_1)$
  $R(u_0, u_1) \to (R(u_0, u_1) \wedge R(u_1, u_2)) \to (R(u_0, u_1) \wedge R(u_1, u_2) \wedge R(u_2, u_3)) \to \cdots$

- **Fact**: if all TGDs are *full* (i.e. no $\exists$) then any chase terminates.

---

[1]The book doesn't consider constants; need to add this to allow constants.

When the Chase Goes Wrong

- Given a set Σ of GD, we can repeatedly apply the chase:
  $Q \overset{\sigma_1, \theta_1}{\to} Q_1 \overset{\sigma_2, \theta_2}{\to} Q_2 \cdots$

- In general, this may not terminate:
  $\sigma = \forall x \forall y (R(x, y) \to \exists z R(y, z))$ $Q() = R(u_0, u_1)$
  $R(u_0, u_1) \to (R(u_0, u_1) \wedge R(u_1, u_2)) \to (R(u_0, u_1) \wedge R(u_1, u_2) \wedge R(u_2, u_3)) \to \cdots$

- **Fact**: if all TGDs are *full* (i.e. no $\exists$) then any chase terminates.

- In general, the chase may also fail:
  $\sigma = \forall x \forall y \forall z (R(x, y) \wedge R(x, z) \Rightarrow (y = z))$ $Q() = R(x, 33) \wedge R(x, 55)$
  $Q \to \texttt{fail}$.

---

[1]The book doesn't consider constants; need to add this to allow constants.

When the Chase Goes Wrong

- Given a set Σ of GD, we can repeatedly apply the chase:
  $Q \xrightarrow{\sigma_1, \theta_1} Q_1 \xrightarrow{\sigma_2, \theta_2} Q_2 \cdots$

- In general, this may not terminate:
  $\sigma = \forall x \forall y (R(x, y) \rightarrow \exists z R(y, z))$ $Q() = R(u_0, u_1)$
  $R(u_0, u_1) \rightarrow (R(u_0, u_1) \wedge R(u_1, u_2)) \rightarrow (R(u_0, u_1) \wedge R(u_1, u_2) \wedge R(u_2, u_3)) \rightarrow \cdots$

- **Fact**: if all TGDs are *full* (i.e. no $\exists$) then any chase terminates.

- In general, the chase may also fail:
  $\sigma = \forall x \forall y \forall z (R(x, y) \wedge R(x, z) \Rightarrow (y = z))$ $Q() = R(x, 33) \wedge R(x, 55)$
  $Q \rightarrow \texttt{fail}$.

- **Fact**: if Σ does not contain EGDs, then the chase never fails.

---

[1] The book doesn't consider constants; need to add this to allow constants.

When the Chase Goes Wrong

- Given a set Σ of GD, we can repeatedly apply the chase:
  $Q \overset{\sigma_1, \theta_1}{\to} Q_1 \overset{\sigma_2, \theta_2}{\to} Q_2 \cdots$

- In general, this may not terminate:
  $\sigma = \forall x \forall y (R(x, y) \to \exists z R(y, z))$ $Q() = R(u_0, u_1)$
  $R(u_0, u_1) \to (R(u_0, u_1) \wedge R(u_1, u_2)) \to (R(u_0, u_1) \wedge R(u_1, u_2) \wedge R(u_2, u_3)) \to \cdots$

- **Fact**: if all TGDs are *full* (i.e. no $\exists$) then any chase terminates.

- In general, the chase may also fail:
  $\sigma = \forall x \forall y \forall z (R(x, y) \wedge R(x, z) \Rightarrow (y = z))$ $Q() = R(x, 33) \wedge R(x, 55)$
  $Q \to \texttt{fail}$.

- **Fact**: if Σ does not contain EGDs, then the chase never fails.

- **Theorem** [Abiteboul et al., 1995, Theorem 8.4.18]: if Σ consists of full TGDs and EDGs (i.e. no $\exists$) and the chase succeeds[1] then all terminating chases end in the same query, denoted $\texttt{Chase}(Q)$.

---

[1]The book doesn't consider constants; need to add this to allow constants.

Soundness

Theorem (Soundness Theorem)

Let $\sigma = \boxed{\forall \boldsymbol{x}(A \Rightarrow C)}$ be a GD. If $Q \overset{\sigma,\theta}{\to} Q_1$ then $\sigma \models Q \subseteq Q_1$.

Recap
00

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
000000●00000

Soundness

Theorem (Soundness Theorem)

Let $\sigma = \boxed{\forall \boldsymbol{x}(A \Rightarrow C)}$ be a GD. If $Q \xrightarrow{\sigma, \theta} Q_1$ then $\sigma \models Q \subseteq Q_1$.

Proof Assume Q is a Boolean query. Let $\boldsymbol{D}$ be s.t. $\boldsymbol{D} \models \sigma$, $Q(\boldsymbol{D}) = \text{true}$.

Soundness

Theorem (Soundness Theorem)

Let $\sigma = \boxed{\forall \boldsymbol{x}(A \Rightarrow C)}$ be a GD. If $Q \overset{\sigma, \theta}{\rightarrow} Q_1$ then $\sigma \models Q \subseteq Q_1$.

Proof Assume Q is a Boolean query. Let $\boldsymbol{D}$ be s.t. $\boldsymbol{D} \models \sigma$, $Q(\boldsymbol{D}) = \text{true}$.

Then $\exists \varphi : Q \to \boldsymbol{D}$.   Compose it with θ: $\varphi \circ \theta = \ A \overset{\theta}{\rightarrow} Q \overset{\varphi}{\rightarrow} \boldsymbol{D}$

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○●○○○○○

Soundness

Theorem (Soundness Theorem)

Let $\sigma = \boxed{\forall \boldsymbol{x}(A \Rightarrow C)}$ be a GD. If $Q \overset{\sigma, \theta}{\to} Q_1$ then $\sigma \models Q \subseteq Q_1$.

Proof Assume Q is a Boolean query. Let $\boldsymbol{D}$ be s.t. $\boldsymbol{D} \models \sigma$, $Q(\boldsymbol{D}) = $ true.

Then $\exists \varphi : Q \to \boldsymbol{D}$.   Compose it with θ: $\varphi \circ \theta = \quad A \overset{\theta}{\to} Q \overset{\varphi}{\to} \boldsymbol{D}$

Case 1: σ is a TGD $\boxed{\forall \boldsymbol{x}(A \Rightarrow \exists \boldsymbol{y} B)}$ Then $Q_1 \overset{\text{def}}{=} Q \wedge \theta(B)$.

Recap
oo

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
000000●00000

Soundness

> ### Theorem (Soundness Theorem)
>
> Let $\sigma = \boxed{\forall \mathbf{x}(A \Rightarrow C)}$ be a GD. If $Q \overset{\sigma,\theta}{\to} Q_1$ then $\sigma \models Q \subseteq Q_1$.

Proof Assume Q is a Boolean query. Let $\mathbf{D}$ be s.t. $\mathbf{D} \models \sigma$, $Q(\mathbf{D}) = $ true.

Then $\exists \varphi : Q \to \mathbf{D}$.   Compose it with θ: $\varphi \circ \theta = \ A \overset{\theta}{\to} Q \overset{\varphi}{\to} \mathbf{D}$

Case 1: σ is a TGD $\boxed{\forall \mathbf{x}(A \Rightarrow \exists \mathbf{y} B)}$ Then $Q_1 \overset{\text{def}}{=} Q \wedge \theta(B)$.
$\varphi \circ \theta(A) \subseteq \mathbf{D}$ and $\mathbf{D} \models \sigma$ implies $\varphi \circ \theta$ extends to a homomorphism
$B \to \mathbf{D}$ that factors as $B \overset{\theta}{\to} Q_1 \to \mathbf{D}$, thus $Q_1(\mathbf{D}) = $ true.

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○●○○○○○

Soundness

Theorem (Soundness Theorem)

Let $\sigma = \boxed{\forall \boldsymbol{x}(A \Rightarrow C)}$ be a GD. If $Q \stackrel{\sigma,\theta}{\to} Q_1$ then $\sigma \models Q \subseteq Q_1$.

Proof Assume Q is a Boolean query. Let $\boldsymbol{D}$ be s.t. $\boldsymbol{D} \models \sigma$, $Q(\boldsymbol{D}) =$ true.

Then $\exists \varphi : Q \to \boldsymbol{D}$.   Compose it with θ: $\varphi \circ \theta = \quad A \stackrel{\theta}{\to} Q \stackrel{\varphi}{\to} \boldsymbol{D}$

Case 1: σ is a TGD $\boxed{\forall \boldsymbol{x}(A \Rightarrow \exists \boldsymbol{y} B)}$ Then $Q_1 \stackrel{\text{def}}{=} Q \wedge \theta(B)$.
$\varphi \circ \theta(A) \subseteq \boldsymbol{D}$ and $\boldsymbol{D} \models \sigma$ implies $\varphi \circ \theta$ extends to a homomorphism
$B \to \boldsymbol{D}$ that factors as $B \stackrel{\theta}{\to} Q_1 \to \boldsymbol{D}$, thus $Q_1(\boldsymbol{D}) =$ true.

Case 2: is an EGD $\boxed{\forall \boldsymbol{x}(A \Rightarrow (x_i = x_j))}$ In class.

Recap
00

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
000000000000

Chase for Query Containment

We want to check $\Sigma \models Q \subseteq Q'$

- Simple (but important) observation. If $Q \rightarrow Q_1$ then $Q_1 \subseteq Q$ (unconditioned). Why?

- The Soundness Theorem proves $\Sigma \models Q \subseteq Q_1$.

- To check $\Sigma \models Q \subseteq Q'$, repeatedly chase Q:
$$Q \rightarrow Q_1 \rightarrow Q_2 \rightarrow \cdots$$
$$\cdots \subseteq Q_2 \subseteq Q_1 \subseteq Q \text{ (unconditioned)}$$

- If $Q_m \subseteq Q'$ (unconditioned) for some $m \geq 0$, then $\Sigma \models Q \subseteq Q'$ why?.

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○●○○○

Chase for Query Equivalence

To check equivalence $\Sigma \models Q \equiv Q'$, we need to chase both Q and Q':

$$Q \to Q_1 \to Q_2 \to \cdots \qquad Q' \to Q'_1 \to Q'_2 \to \cdots$$

If $Q_m \equiv Q'_n$ for some m, n, then $\Sigma \models Q \equiv Q'$

Chase and Backchase

[Popa et al., 2000]
Semantics optimization of Q under constraints Σ.

Assume Σ has only _full_ TGDs and EGDs.

Chase Chase Q to completion: $Q \overset{*}{\to} \mathtt{Chase}(Q)$.

Backchase Go in reverse $\mathtt{Chase}(Q) \leftarrow Q_1' \leftarrow Q_2' \leftarrow \cdots$
There are multiple choices for the backchase: this is an optimization
problem.

Recap
○○

Classical Database Constraints
○○○○○○○

Generalized Dependencies
○○○○

Semantics Optimization
○○○○○○○○○○●○

Example: Using Physical Access Structures

Relation $R(k, x, y)$, key k, index $I(k, x)$ on $R.x$

Want to optimize $\boxed{Q(y) = R(k, 55, y)}$ to $\boxed{Q'(y) = R(k, x, y) \wedge I(k, 55)}$

Example: Using Physical Access Structures

Relation $R(k, x, y)$, key k, index $I(k, x)$ on $R.x$

Want to optimize $\boxed{Q(y) = R(k, 55, y)}$ to $\boxed{Q'(y) = R(k, x, y) \wedge I(k, 55)}$

> FD σ_0: $\forall k, x_1, x_2, y_1, y_2(R(k, x_1, y_1) \wedge R(k, x_2, y_2) \Rightarrow (x_1 = x_2))$
> IND1: σ_1: $\forall k, x, y(R(k, x, y) \Rightarrow I(k, x))$
> IND2: σ_2: $\forall k I(k, x) \rightarrow \exists y R(k, x, y)$

Recap
oo

Classical Database Constraints
oooooooo

Generalized Dependencies
oooo

Semantics Optimization
ooooooooooooo

Example: Using Physical Access Structures

Relation $R(k, x, y)$, key k, index $I(k, x)$ on $R.x$

Want to optimize $\boxed{Q(y) = R(k, 55, y)}$ to $\boxed{Q'(y) = R(k, x, y) \wedge I(k, 55)}$

FD σ_0: $\forall k, x_1, x_2, y_1, y_2 (R(k, x_1, y_1) \wedge R(k, x_2, y_2) \Rightarrow (x_1 = x_2))$

IND1: σ_1: $\forall k, x, y(R(k, x, y) \Rightarrow I(k, x))$

IND2: σ_2: $\forall k I(k, x) \rightarrow \exists y R(k, x, y)$

$$
\begin{array}{l}
Q \equiv \quad R(k, 55, y) \quad \overset{\sigma_1}{\rightarrow} \quad R(k, 55, y) \wedge I(k, 55) \quad \equiv \texttt{Chase}(Q). \\
Q' \equiv R(k, x, y) \wedge I(k, 55) \overset{\sigma_2}{\rightarrow} R(k, x, y) \wedge R(k, 55, y') \wedge I(k, 55) \\
\qquad \overset{\sigma_0}{\rightarrow} R(k, 55, y) \wedge I(k, 55) \qquad\qquad \equiv \texttt{Chase}(Q')
\end{array}
$$

$\texttt{Chase}(Q) = \texttt{Chase}(Q')$, implies $\Sigma \models Q \equiv Q'$.

Recap
00

Classical Database Constraints
0000000

Generalized Dependencies
0000

Semantics Optimization
000000000000

Example: Using Physical Access Structures

Relation $R(k, x, y)$, key k, index $I(k, x)$ on $R.x$

Want to optimize $\boxed{Q(y) = R(k, 55, y)}$ to $\boxed{Q'(y) = R(k, x, y) \wedge I(k, 55)}$

FD σ_0: $\forall k, x_1, x_2, y_1, y_2(R(k, x_1, y_1) \wedge R(k, x_2, y_2) \Rightarrow (x_1 = x_2))$
IND1: σ_1: $\forall k, x, y(R(k, x, y) \Rightarrow I(k, x))$
IND2: σ_2: $\forall k I(k, x) \rightarrow \exists y R(k, x, y)$

$$
\begin{array}{rl}
Q \equiv & R(k, 55, y) \quad \overset{\sigma_1}{\rightarrow} \quad R(k, 55, y) \wedge I(k, 55) \quad \equiv \texttt{Chase}(Q). \\
Q' \equiv & R(k, x, y) \wedge I(k, 55) \overset{\sigma_2}{\rightarrow} R(k, x, y) \wedge R(k, 55, y') \wedge I(k, 55) \\
& \overset{\sigma_0}{\rightarrow} R(k, 55, y) \wedge I(k, 55) \qquad\qquad \equiv \texttt{Chase}(Q')
\end{array}
$$

$\texttt{Chase}(Q) = \texttt{Chase}(Q')$, implies $\Sigma \models Q \equiv Q'$.

Given Q, chase/Backchase computes $\texttt{Chase}(Q)$ the <u>searches</u> for Q':

$$Q \overset{\sigma_1}{\rightarrow} \quad \overset{\sigma_0}{\leftarrow} \quad \overset{\sigma_2}{\leftarrow} Q'$$

Summary

- Constraints are restricted sentences in FO.

- The implication problem: elegant theory because it's a special case of logical implication.

- Chase: simple, yet fundamental technique. Egg and Egglog use chase. Termination is undecidable in general. Understanding termination is a major open problem.

Abiteboul, S., Hull, R., and Vianu, V. (1995).
Foundations of Databases.
Addison-Wesley.

Chakravarthy, U. S., Grant, J., and Minker, J. (1990).
Logic-based approach to semantic query optimization.
ACM Trans. Database Syst., 15(2):162–207.

King, J. J. (1981).
QUIST: A system for semantic query optimization in relational databases.
In *Very Large Data Bases, 7th International Conference, September 9-11, 1981, Cannes, France, Proceedings*, pages 510–517. IEEE Computer Society.

Popa, L., Deutsch, A., Sahuguet, A., and Tannen, V. (2000).
A chase too far?
In Chen, W., Naughton, J. F., and Bernstein, P. A., editors, *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data, May 16-18, 2000, Dallas, Texas, USA*, pages 273–284. ACM.